

Ecological forecasting with dynamic GAMs

Nicholas Clark

School of Veterinary Science, University of Queensland, Australia

1400 Thursday 6th November, 2023

“Because all decision making is based on what will happen in the future, either under the status quo or different decision alternatives, decision making ultimately depends on forecasts”

Dietze et al. 2018



Properties of ecological series

Temporal autocorrelation

Lagged effects

Non-Gaussian data and missing observations

Measurement error

Time-varying effects

Nonlinearities

Multi-series clustering

Properties of ecological series

Temporal autocorrelation

Lagged effects

Non-Gaussian data and missing observations

Measurement error


Time-varying effects

Nonlinearities


Multi-series clustering



GAMs use splines...



...penalized to fit data



Easy to fit in


$$\mathbb{E}(\mathbf{Y}_t | \mathbf{X}_t) = g^{-1}\left(\alpha + \sum_{j=1}^J f(x_{jt})\right)$$


Where:

g^{-1} is the inverse of the link function


α is the intercept

$f(x)$ are potentially nonlinear functions of the J predictors







Basis: GP



Basis: Tensor product



Basis: TPRS



Basis: CRS (shrink)

Need more on GAMs?

• [GAMs in R](#)

• [GAMs in R](#)

• [GAMs in R](#)

• [GAMs in R](#)

• [GAMs in R](#)

Gavin has you covered 😊

Generalized Additive Models with R and mgcv

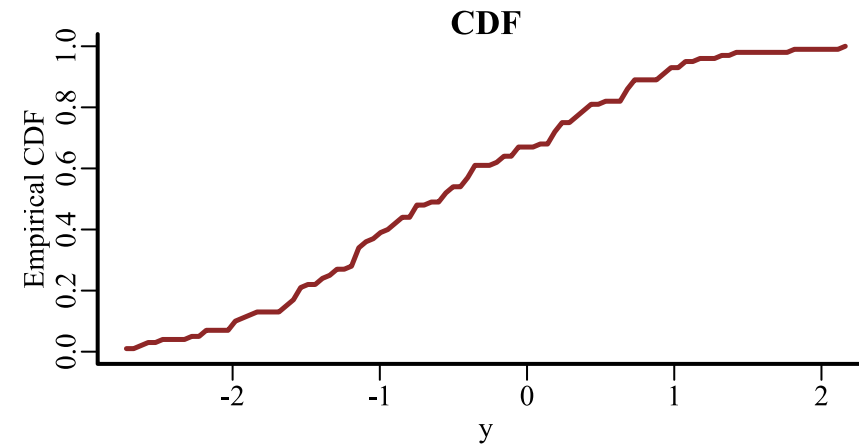



Gavin Simpson

January 3, 2022

GAMs are just fancy GLMs, where some (or all) of the predictor effects are estimated as (possibly nonlinear) smooth functions

**But the complexity these smooth functions can handle is
*enormous***

What's the catch?




A spline of time


```
library(mgcv)
model <- gam(y ~ s(time, k = 20, bs = 'bs', m = 2),
             data = data,
             family = gaussian())
```

A B-spline (`bs = 'bs'`) with `m = 2` sets the penalty on the second derivative

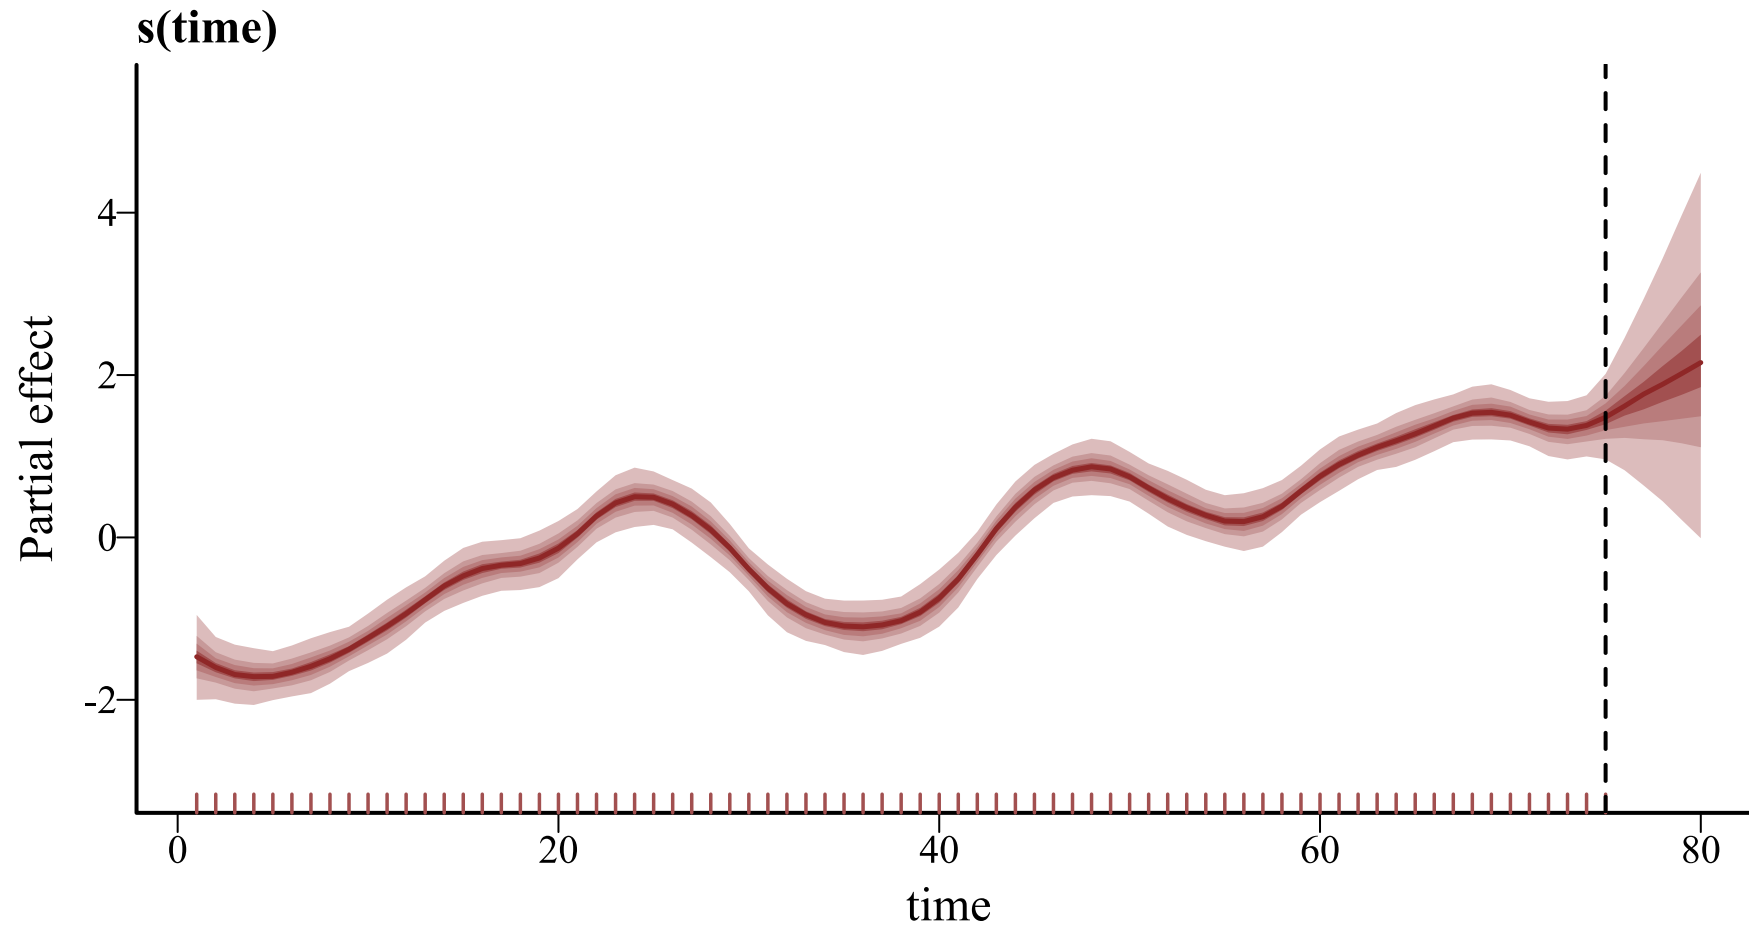
Hindcasts ☺




Extrapolate 2-steps ahead 😊




5-steps ahead ☹️




20-steps ahead ☹️




Basis functions \Rightarrow local knowledge




Basis functions \Rightarrow local knowledge




Basis functions \Rightarrow local knowledge




Basis functions \Rightarrow local knowledge




Forecasts




We need *global knowledge*




We need *global knowledge*



We need *global knowledge*



We need *global knowledge*



Dynamic GAMs

$$\mathbb{E}(\mathbf{Y}_t | \mathbf{X}_t) = g^{-1}(\alpha + \sum_{j=1}^J f(x_{jt}) + z_t)$$

Where:

g^{-1} is the inverse of the link function

α is the intercept

$f(x)$ are potentially nonlinear functions of the J predictors

z_t is a ***latent dynamic process***

Modelling with the mvgam

Bayesian framework to fit Dynamic GLMs and Dynamic GAMs

Hierarchical intercepts, slopes and smooths

Latent dynamic processes

State-Space models with measurement error

Built off the mgcv  to construct penalized smoothing splines

Convenient and familiar  formula interface

Uni- or multivariate series from a range of response distributions

Uses Stan for efficient Hamiltonian Monte Carlo sampling

Observation families

`gaussian()`, `student-t()` \Rightarrow real values in $(-\infty, \infty)$

`lognormal()`, `Gamma()` \Rightarrow positive real values in $[0, \infty)$

`betar()` \Rightarrow real values (proportional) in $[0, 1]$

`poisson()`, `nb()` \Rightarrow non-negative integers in $(0, 1, 2, \dots)$

Extended predictor effects

`s()` \Rightarrow Smoothing spline of one or more covariates

`s(bs = 're')` \Rightarrow Hierarchical slopes or intercepts

`te()`, `ti()`, `t2()` \Rightarrow Tensor product smoothing spline of two or more covariates

`gp()` \Rightarrow Gaussian Process function (with squared exponential kernel) of one covariate


`dynamic()` \Rightarrow Time-varying effect of one covariate

We can fit models that include random effects, nonlinear effects and complex multidimensional smooth functions. All these effects can operate *on both process and observation models*


Can incorporate unobserved temporal dynamics; do not need to regress the outcome on its own past values

Very advantageous for ecological time series. But what kinds of dynamic processes are available in the `mvgam` ?


Random Walk or AR(1-3)




Gaussian Process...




...where length scale \Rightarrow *memory*



VAR1 \Rightarrow Granger causality



Factors \Rightarrow induced correlations



Example of the interface

```
model ← mvgam(  
  formula = y ~  
    s(series, bs = 're') +  
    s(x0, series, bs = 're') +  
    x1 +  
    gp() +  
    te(x3, x4, bs = c('cr', 'tp')),  
  data = data,  
  family = poisson(),  
  trend_model = 'AR1',  
  burnin = 500,  
  samples = 500,  
  chains = 4,  
  parallel = TRUE  
)
```

Example data (long format)

| y | series | time |
|----|-----------|------|
| 2 | species_1 | 1 |
| 0 | species_2 | 1 |
| NA | species_3 | 1 |
| NA | species_4 | 1 |
| 1 | species_1 | 2 |
| 0 | species_2 | 2 |
| 3 | species_3 | 2 |
| 5 | species_4 | 2 |

Response (NAs allowed)

| y | series | time |
|----|-----------|------|
| 2 | species_1 | 1 |
| 0 | species_2 | 1 |
| NA | species_3 | 1 |
| NA | species_4 | 1 |
| 1 | species_1 | 2 |
| 0 | species_2 | 2 |
| 3 | species_3 | 2 |
| 5 | species_4 | 2 |

Series indicator (as factor)

| y | series | time |
|----|-----------|------|
| 2 | species_1 | 1 |
| 0 | species_2 | 1 |
| NA | species_3 | 1 |
| NA | species_4 | 1 |
| 1 | species_1 | 2 |
| 0 | species_2 | 2 |
| 3 | species_3 | 2 |
| 5 | species_4 | 2 |


Time indicator

| y | series | time |
|----|-----------|------|
| 2 | species_1 | 1 |
| 0 | species_2 | 1 |
| NA | species_3 | 1 |
| NA | species_4 | 1 |
| 1 | species_1 | 2 |
| 0 | species_2 | 2 |
| 3 | species_3 | 2 |
| 5 | species_4 | 2 |

Any other predictors

| y | series | time | x0 | x1 | x2 | x3 | x4 |
|----|-----------|------|-------|----|-------|-------|-------|
| 2 | species_1 | 1 | -0.38 | A | 0.20 | 1.18 | -0.72 |
| 0 | species_2 | 1 | -0.71 | A | -2.67 | 1.02 | 0.67 |
| NA | species_3 | 1 | 0.05 | B | -0.33 | 0.12 | 1.50 |
| NA | species_4 | 1 | 0.77 | B | 0.65 | 0.86 | -0.49 |
| 1 | species_1 | 2 | 0.29 | A | -0.25 | 1.18 | -0.82 |
| 0 | species_2 | 2 | 0.34 | A | -0.15 | 2.12 | 0.20 |
| 3 | species_3 | 2 | -0.38 | B | -0.81 | 1.33 | -1.15 |
| 5 | species_4 | 2 | 1.32 | B | 0.22 | -0.72 | 1.36 |

Types of `mvgam` predictions




modified from [Heiss 2022](#)

Workflow in `mvgam`

Fit models that can include nonlinear splines, GPs, and multivariate dynamic processes to ecological time series

Use posterior predictive checks and Randomized Quantile (Dunn-Smyth) residuals to assess model failures

Use `marginaleffects`  to generate interpretable (and reportable) model predictions

Produce probabilistic forecasts

Evaluate forecasts from competing models with proper scoring rules

More resources

Vignette ⇒ [Overview of the package](#)

Vignette ⇒ [Formatting data for use in `mvgam`](#)

Vignette ⇒ [Shared latent process models](#)

Vignette ⇒ [Time-varying effects](#)

Vignette ⇒ [Multivariate State-Space models](#)

Motivating publication ⇒ Clark & Wells 2023 [*Methods in Ecology and Evolution*](#)